



The IEEE/ACM International Conference on
Software Engineering (ICSE 2025)



LiCoEval: Evaluating LLMs on License Compliance in Code Generation



Weiwei Xu

xuww@stu.pku.edu.cn


Peking University



Kai Gao

kai.gao@ustb.edu.cn

University of Science and Technology Beijing



Hao He

haohe@andrew.cmu.edu

Carnegie Mellon University



Minghui Zhou

zhmh@pku.edu.cn

Peking University

Ottawa, May 1, 2025

AI coding tools have been widely adopted but raised growing controversy about copyright

92% Developers

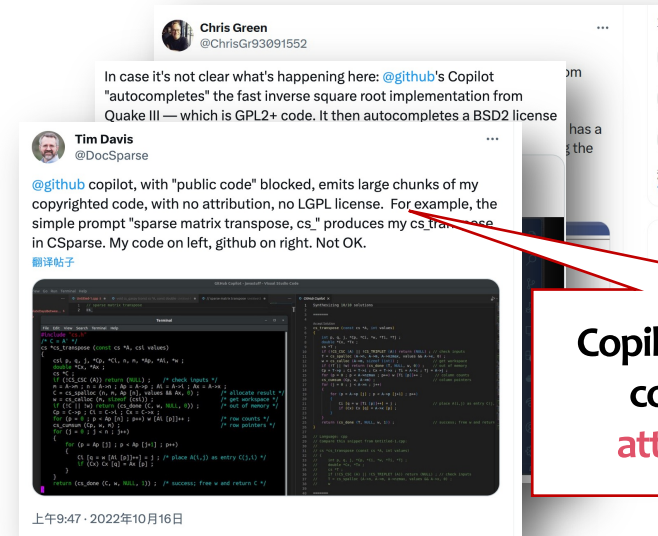
In U.S.

Using AI coding Tools

1/3 Projects

With at least one star

Using GitHub Copilot



Copilot emits large chunks of my copyrighted code, **with no attribution, no LGPL license.**

Social Media Posts

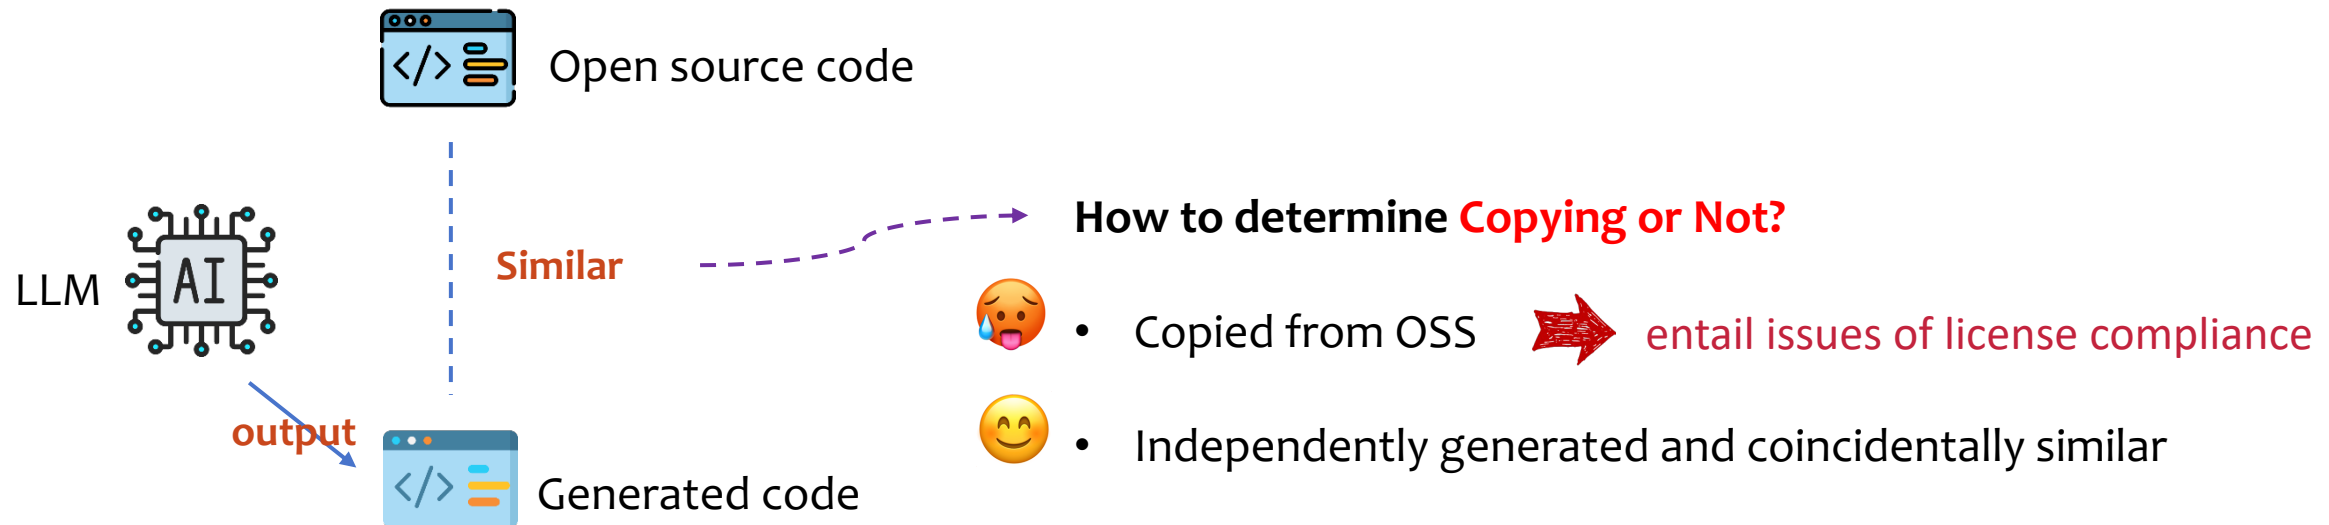
VII. FACTUAL ALLEGATIONS.....

We've filed a lawsuit challenging GitHub Copilot, an AI product that relies on unprecedent open-source software piracy. **Because AI needs to be legal & ethical for everyone**

Copilot Outputs Copyrighted Materials Without Following the Terms of the Applicable License

Evaluating LLM's license compliance in code generation is important but challenging

- Evaluate their ability to **provide accurate license and copyright information** during code generation
 - protect the IP rights of numerous open-source developers
 - shield users of such LLMs from unforeseen legal risks
- Challenge: **Copying or coincidentally similar?** 🤔



The key principle in law determining copyright infringement comes to the rescue

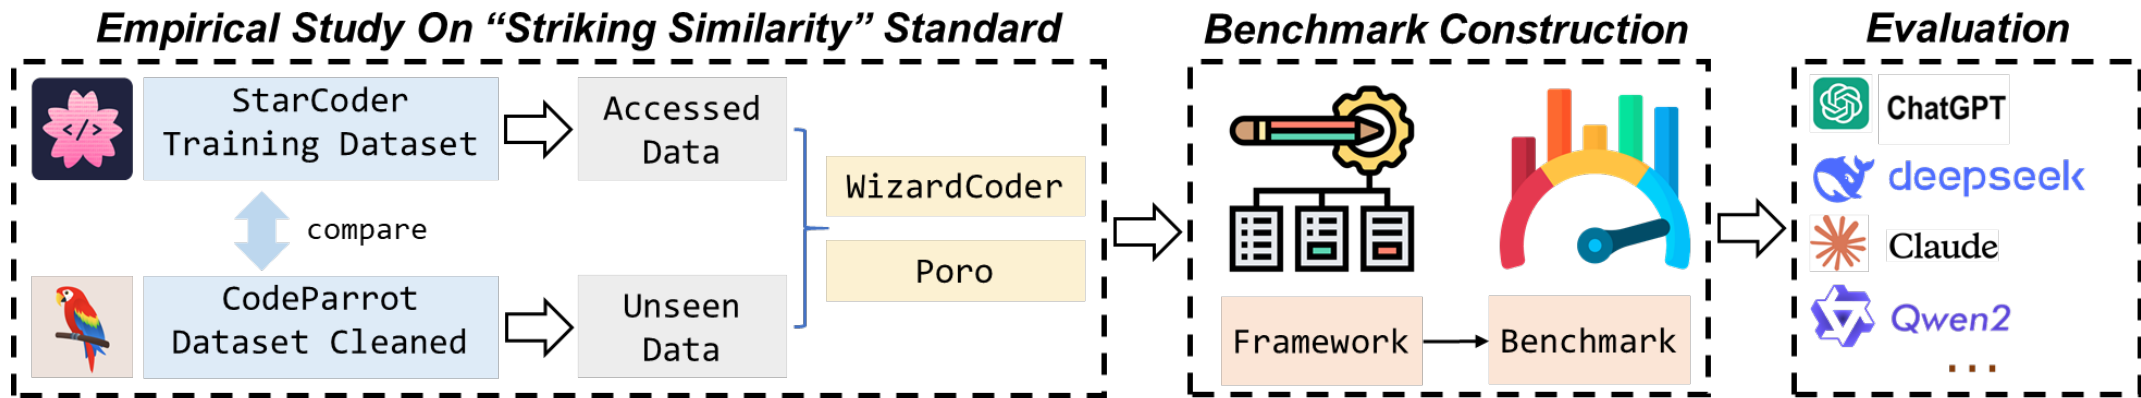
Courts have identified **two methods** to prove substantial copying of an original work^[1].

- The plaintiff can choose to provide evidence showing that the defendant had “**access**” to the copyrighted work and that the two works are “**substantially similar**.” ❌ Because LLM’s training data is usually undisclosed, “access” can not be proven.
- On the other hand, when access cannot be proven, the plaintiff can provide evidence demonstrating that the works in question have “**striking similarity**” (sufficient to rule out the possibility of independent creation). ✅



[1] U. S. C. for the Ninth Circuit. Copying—access and substantial similarity. [Online]. Available: <https://www.ce9.uscourts.gov/jury-instructions/node/274>

Our Solution: evaluating license compliance of LLMs based on Striking Similarities



Overview of this study

Establishing Striking Similarity Standard by Comparing LLM's Performance on Generating Accessed and Unseen Data

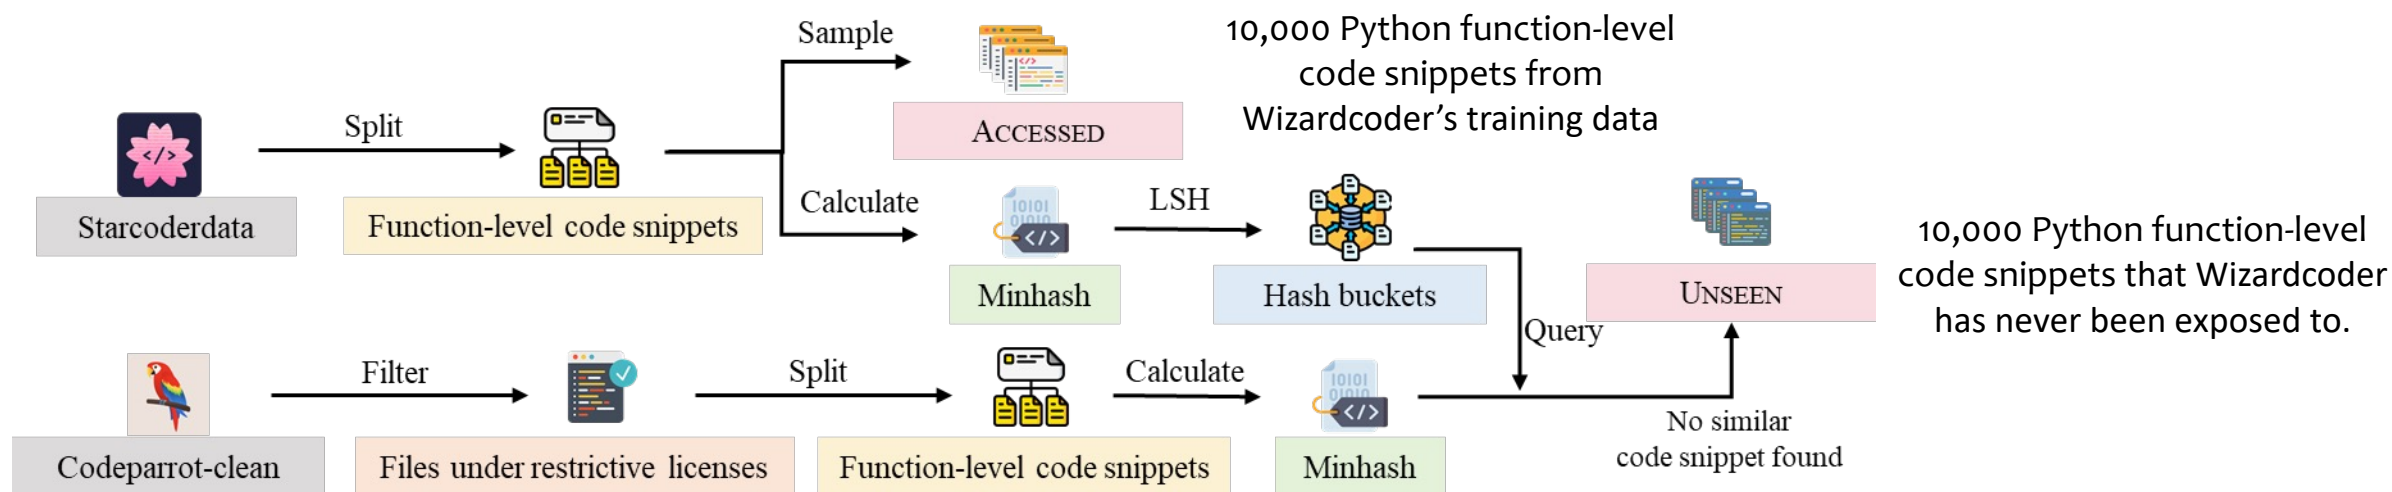
Where might the reasonable standard of striking similarity lie in the context of code generation by LLMs?

Model for analysis: WizardCoder-15B-V1.0

😊 All training data is public.

Experiment setup:

- construct two distinct groups of code samples, **UNSEEN** and **ACCESSED**, to simulate two different scenarios, i.e., **independent creation** and **non-independent creation**.



Our goal is to observe potential differences in similarity when the model generates code for these two distinct groups.

Constructing prompts using function headers

Model for analysis: WizardCoder-15B-V1.0

Experiment setup:

- construct prompts using the UNSEEN and ACCESSED groups, then instruct WizardCoder to complete the code snippets

```
""" Code for unpacking zip files from iLearn """  
import zipfile  
TAR = '/usr/bin/tar'  
  
def unzip(zfile, outdir):  
    """  
    Unpack a zip file into the given output directory outdir  
    Return True if it worked, False otherwise  
    """  
    try:  
        zf = zipfile.ZipFile(zfile)  
        zf.extractall(outdir)  
        return True  
    ... (omitted due to space limitations)
```

Comments in file header

Import statements and global variables

Function signature

Docstring

Function body




prompt

Complete the following Python function:

```
""" Code for unpacking zip files from iLearn """  
import zipfile  
TAR = '/usr/bin/tar'  
def unzip(zfile, outdir):  
    """  
    Unpack a zip file into the given output directory outdir  
    Return True if it worked, False otherwise  
    """
```

Structure of function-level code snippet

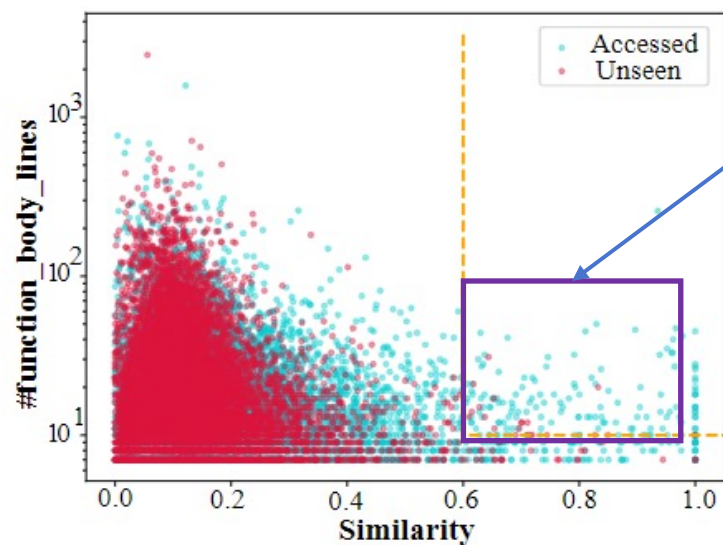
Different performances of WizardCoder for two groups



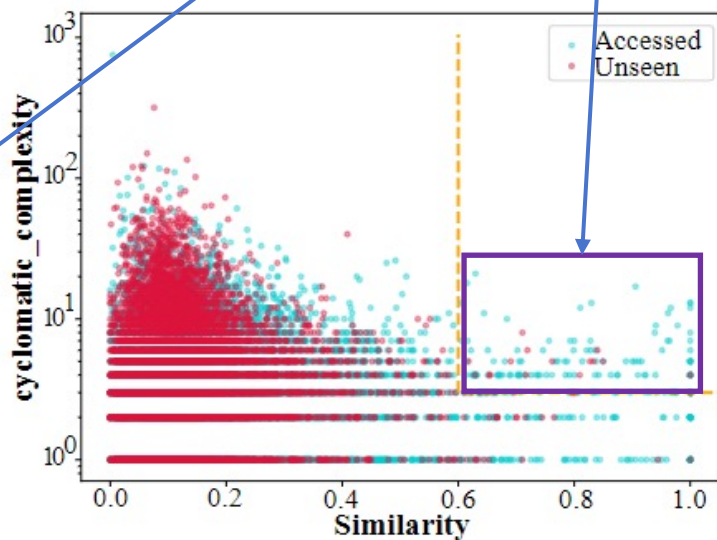
Text similarity alone cannot determine non-independent creation in LLM-generated code.

For complex functions in the UNSEEN group, the similarity is typically lower.

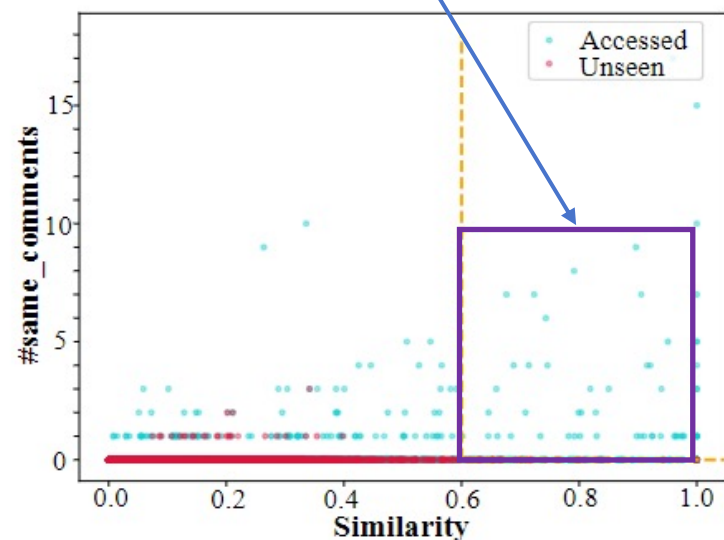
LLMs can memorize and reproduce comments from training data.



(a)



(b)



(c)

A simple but effective standard for “Striking Similarity”

Striking Similarity

number of function body lines > 10

cyclomatic complexity > 3

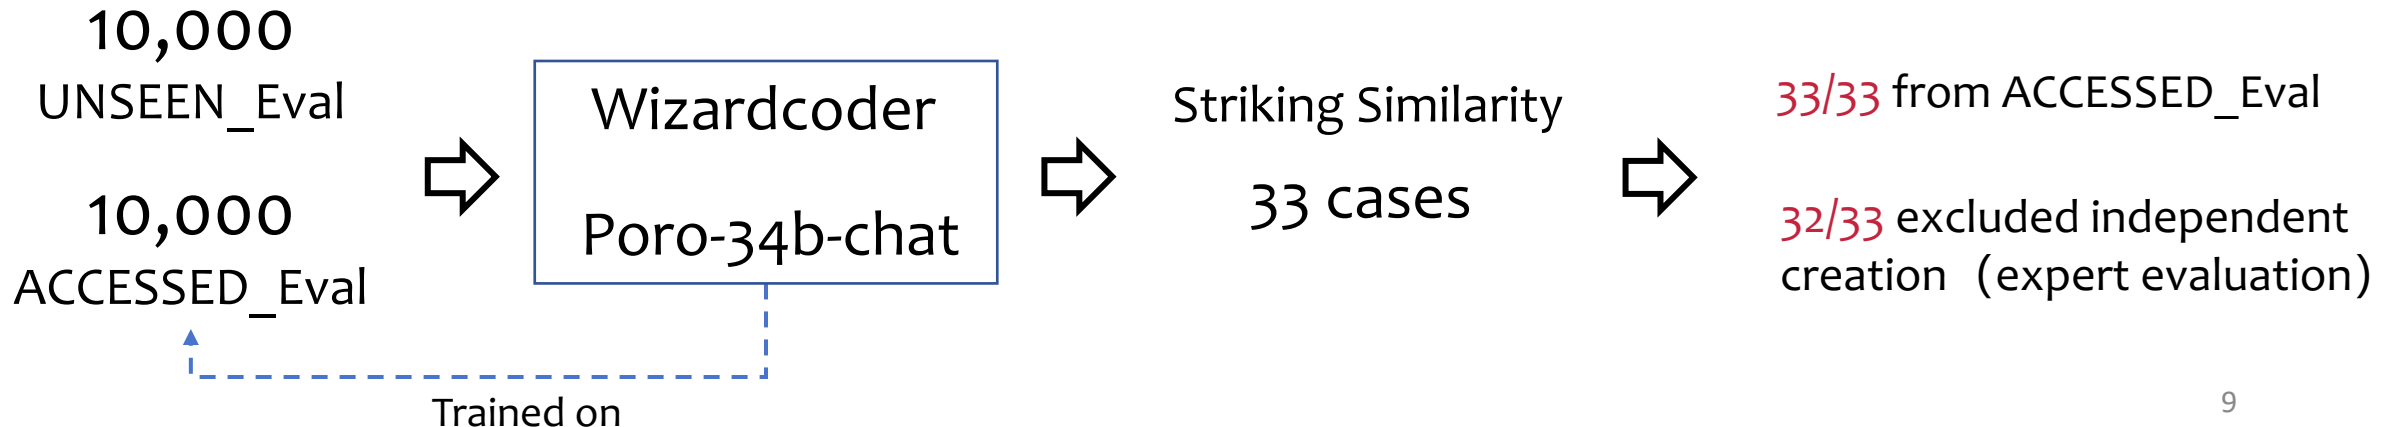
text similarity > 0.6

number of identical comments > 0

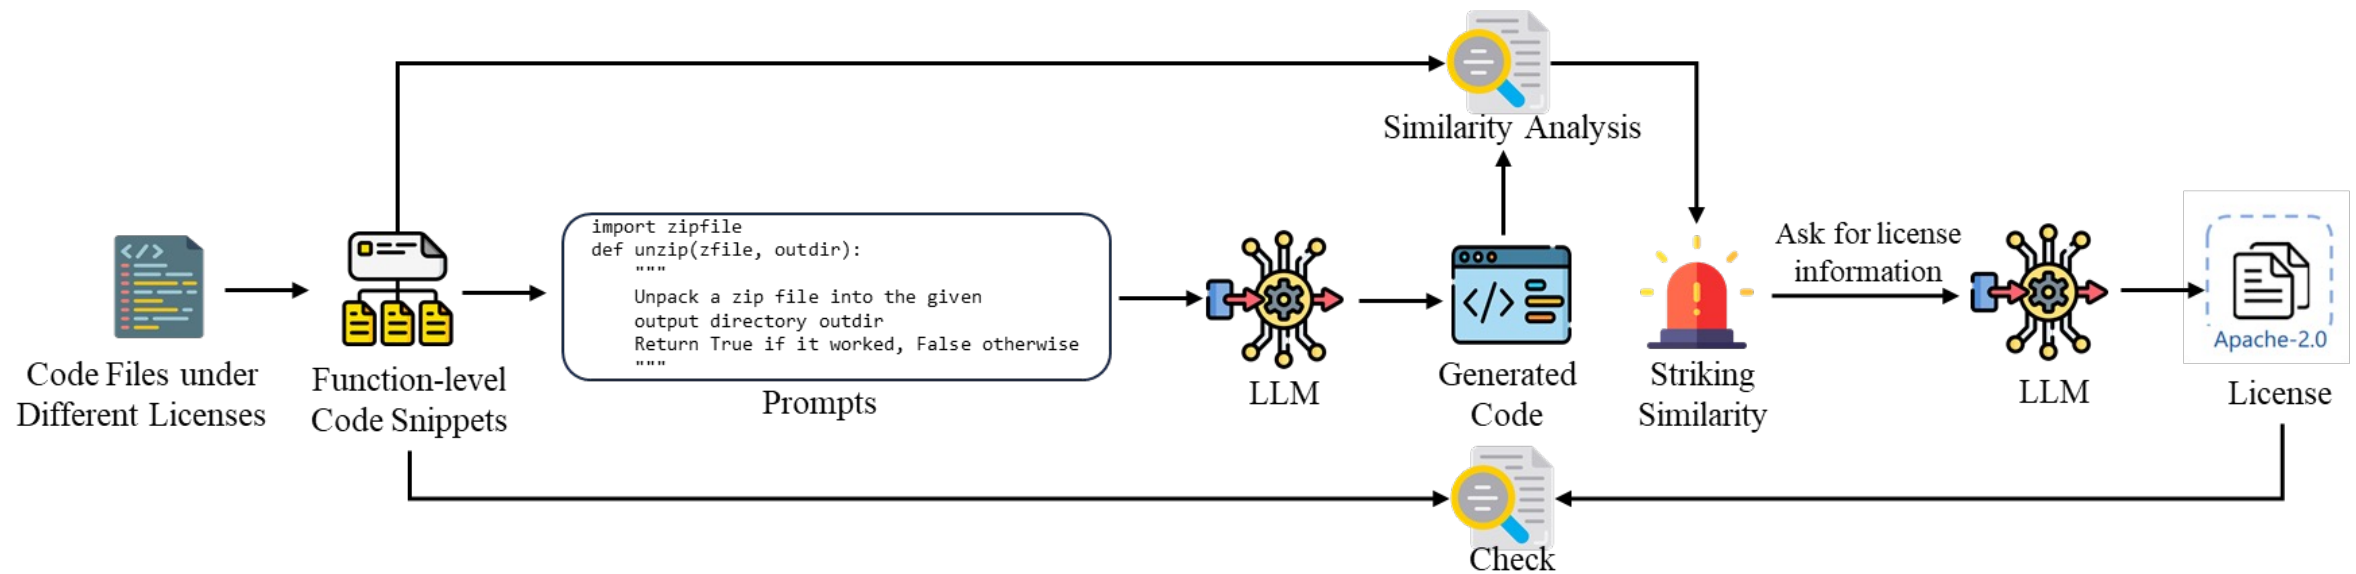


Copyright laws only protect **expression** (i.e., the specific expression of code), **not ideas**.

Evaluation

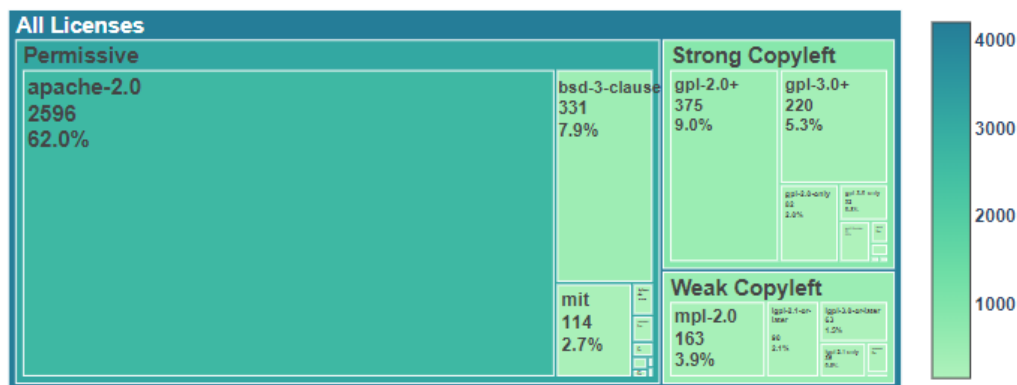
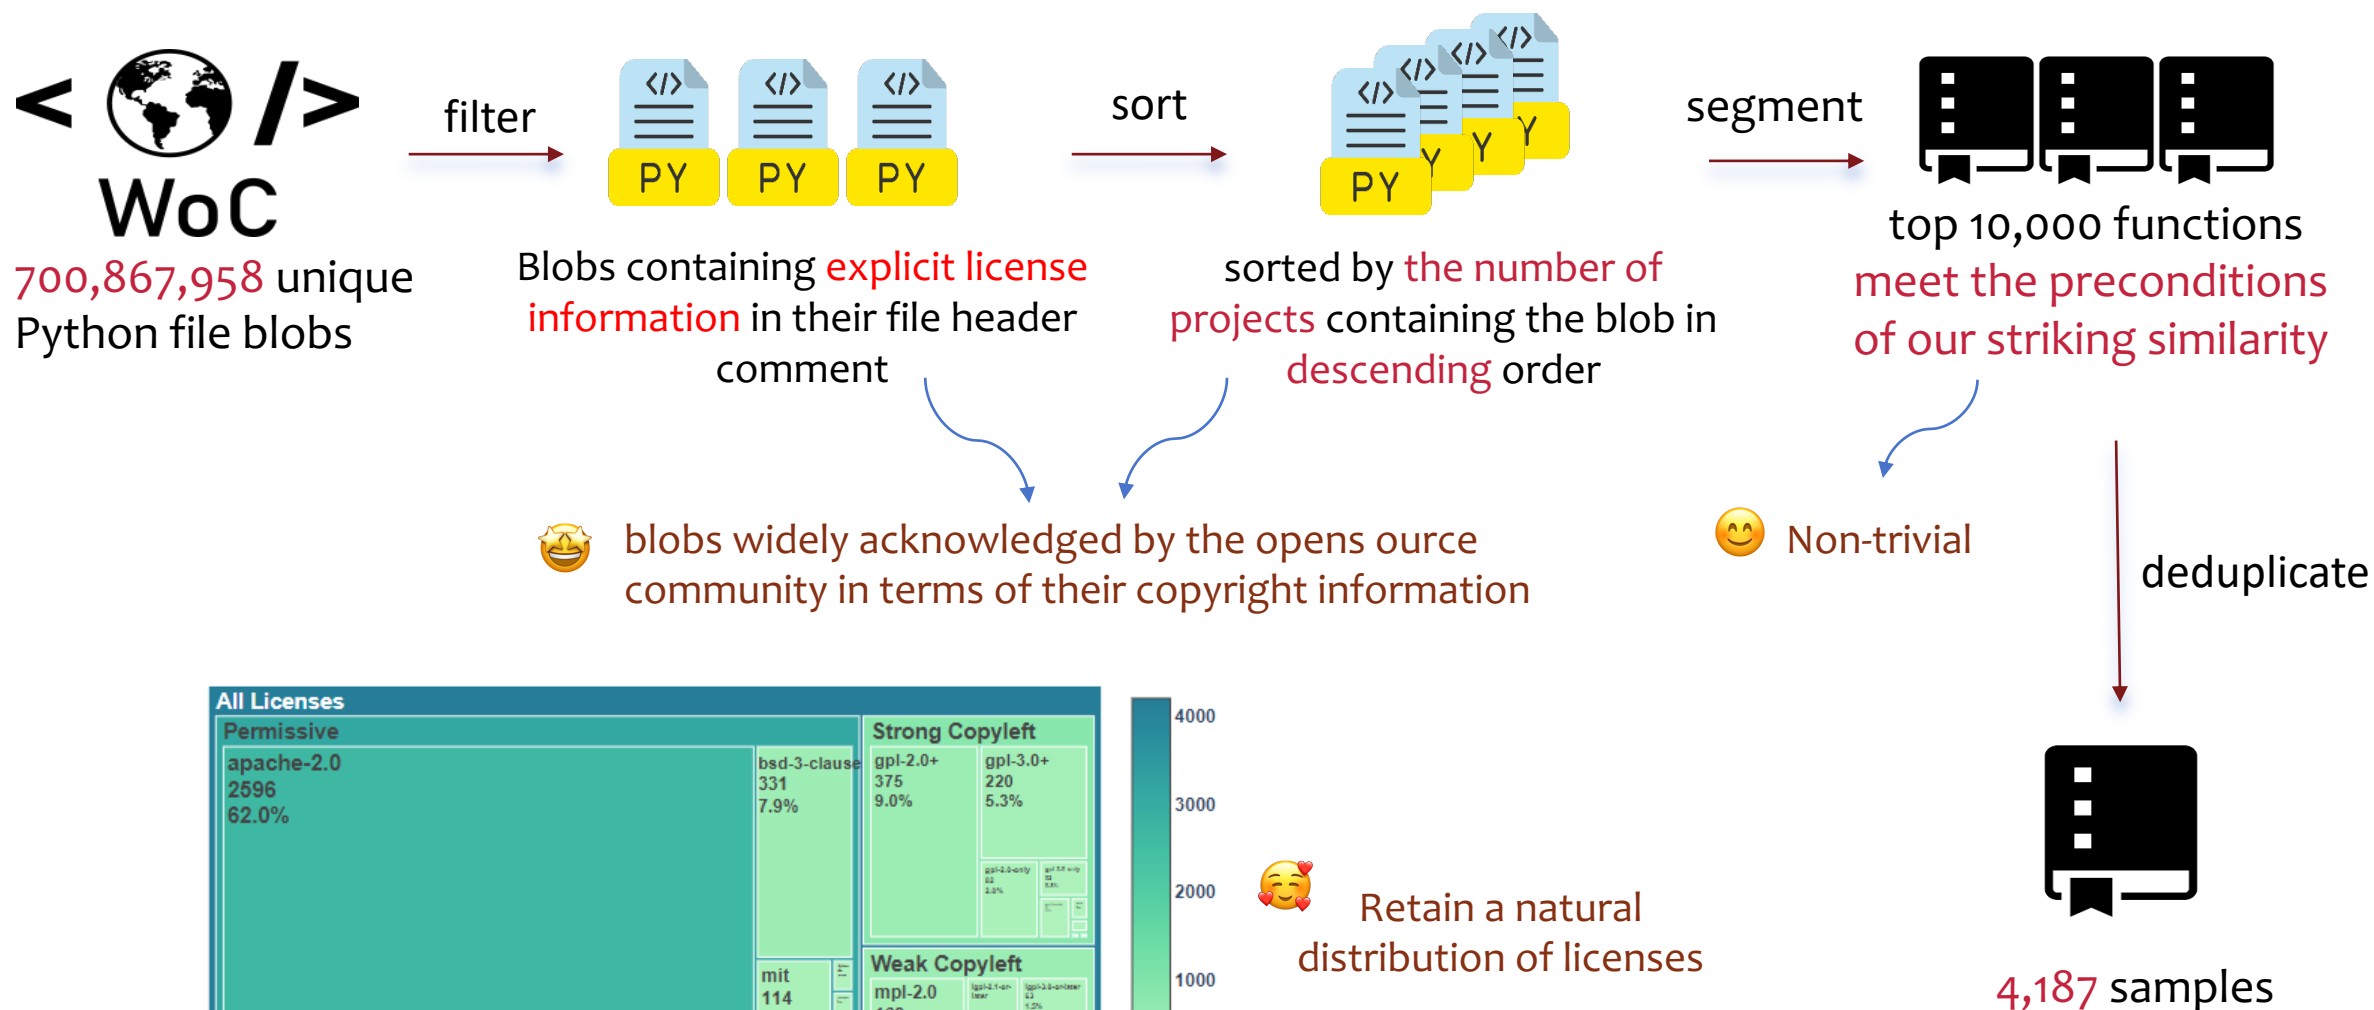


The evaluation framework inspired by the empirical study



Overview of the evaluation framework

Constructing benchmark based on World of Code



The distribution of license in LiCoEval

Evaluating 14 LLMs on LicoEval

PERFORMANCE OF 14 LLMs ON LiCoEVAL. ✓ MEANS PUBLICLY AVAILABLE WEIGHTS AND × MEANS UNAVAILABLE WEIGHTS.

	Model	HumanEval	Weights	#striking_sim	Acc	#permissive	Acc _p	#copyleft	Acc _c
General LLM	GPT-3.5-Turbo	72.6	×	29 (0.69%)	0.72	26	0.81	3	0.0
	GPT-4-Turbo	85.4	×	25 (0.60%)	0.72	22	0.82	3	0.0
	GPT-4o	90.2	×	47 (1.12%)	0.74	41	0.85	6	0.0
	Gemini-1.5-Pro	71.9	×	41 (0.98%)	0.59	39	0.62	2	0.0
	Claude-3.5-Sonnet	92.0	×	84 (2.01%)	0.69	79	0.71	5	0.4
	Qwen2-7B-Instruct	79.9	✓	20 (0.48%)	0.95	20	0.95	0	-
	GLM-4-9B-Chat	71.8	✓	0 (0.0%)	-	-	-	-	-
	Llama-3-8B-Instruct	62.2	✓	1 (0.02%)	0.0	1	0.0	0	-
Code LLM	DeepSeek-Coder-V2	90.2	✓	37 (0.88%)	0.0	36	0.0	1	0.0
	CodeQwen1.5-7B-Chat	83.5	✓	17 (0.41%)	0.24	17	0.24	0	-
	StarCoder2-15B-Instruct	72.6	✓	13 (0.31%)	0.23	13	0.23	0	-
	Codestral-22B-v0.1	61.5	✓	91 (2.17%)	0.73	87	0.77	4	0.0
	CodeGemma-7B-IT	56.1	✓	3 (0.07%)	0.33	3	0.33	0	-
	WizardCoder-Python-13B	64.0	✓	27 (0.64%)	0.04	26	0.04	1	0.0

Discussion

Limitation:

- a “minimum” standard that emphasizes precision and interpretability
- may perform poorly on recall

Even with such a minimum standard, we are still able to obtain concerning results from state-of-the-art LLMs.

What's more:

- Different prompts...
- Different scopes(File? Class?)
- Different languages

Discussion

Implications:

LLM providers:

Data Cleaning and License Detection

Enhancing License-Code Association

Addressing Copyleft Information Suppression

LLM users:

Be aware of risks

Verify generated code

Open-source communities:

Adopting more explicit license declarations

Developing guidelines for incorporating and attributing AI-generated code in projects


Establishing clear policies on how their own code should be used in AI training process

Legal professionals:

It is feasible to characterize non-independent creation in LLM outputs using specific feature.




The IEEE/ACM International Conference on Software Engineering (ICSE 2025)



Paper

Thank you!



Code



Weiwei Xu

xuww@stu.pku.edu.cn


Peking University



Kai Gao

kai.gao@ustb.edu.cn

University of Science and Technology Beijing



Hao He

haohe@andrew.cmu.edu

Carnegie Mellon University



Minghui Zhou

zhmh@pku.edu.cn

Peking University

Acknowledgement: National Natural Science Foundation of China 62332001.